

Machine Learning 1

Dominique Lie (A15470100)

10/21/2021


First up is clustering methods

Kmeans clustering

The function in base R to do Kmeans clustering is called 'kmeans()'

Generate some example data where we know what the answer should be:

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))  
x <- cbind(x=tmp, y=rev(tmp))  
  
plot(x)
```



```
#rnorm generates random data that is normalized
```

Q. Can we use `kmeans()` to cluster the data?

[illegible]

Q. How many points are in each cluster?

 km^{size}

```
## [1] 30 30
```

Q. What ‘component’ of your result object details -cluster size? (refer previous question) -cluster assignment/membership? -cluster center?

```
km$cluster
```


[illegible]

km\$centers

```
##          x          y
## 1 -2.828659  2.973690
## 2  2.973690 -2.828659
```

Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col = km$cluster)
points(km$centers, col = "blue", pch = 15, cex = 2)
```



hclust

A big limitation with kmeans is that we have to tell it K (the number of clusters we want) Analyze this same data with hclust()


Demonstrate the use of dist(), hclust(), plot() , and cutree() functions to do clustering, Generate dendogras and return cluster assignment/membership vector...

```
hc <- hclust(dist(x))
hc
```

```
##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

There is a plot method for hclust result objects. Let's see it.

```
plot(hc)
```



To get our cluster membership vector we have to do a wee bit more work. We have to “cut” the tree where we think it makes sense. For this we use the ‘`cutree()`’ function

```
cutree(hc, h = 6)
```

[illegible]

You can also call `cutree()` setting `k` = the number of grps/clusters you want.


```
cutree(hc, k = 2)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
grps <- cutree(hc, k = 2)
```

Make our results plot

```
plot(x, col = 2)
```



Principal Component Analysis

Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
## [1] 17
```

```
ncol(x)
```

```
## [1] 5
```

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese           105   103      103        66
## Carcass_meat     245   227      242       267
## Other_meat       685   803      750       586
## Fish             147   160      122        93
## Fats_and_oils    193   235      184       209
## Sugars           156   175      147       139
```

Not a great method because rerunning code will keep removing columns

```
dim(x)
```

```
## [1] 17  4
```

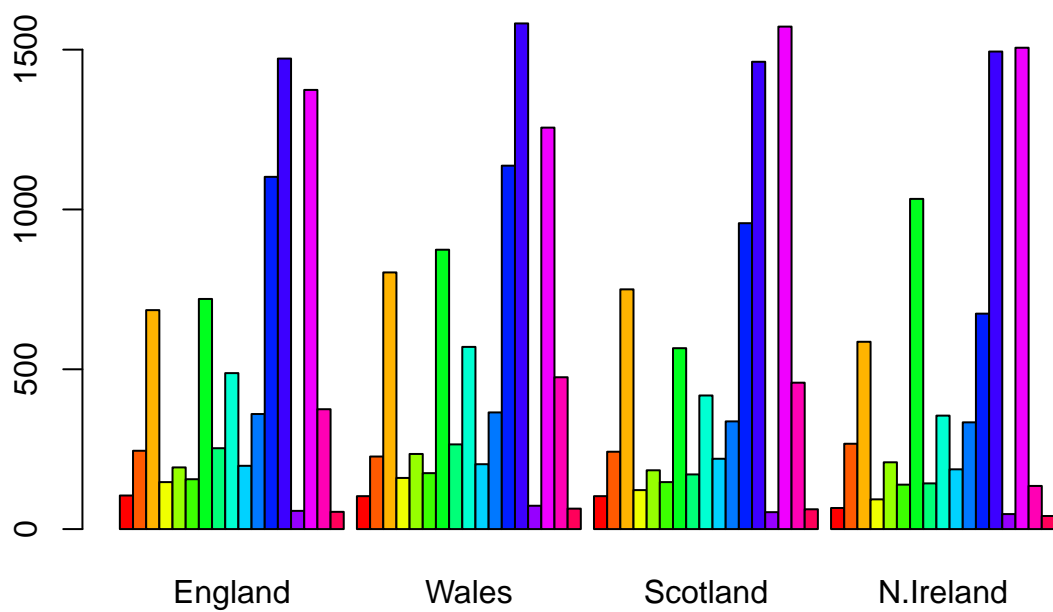
```
read.csv(url, row.names = 1)
```

```
##           England Wales Scotland N.Ireland
## Cheese           105   103      103        66
## Carcass_meat     245   227      242       267
## Other_meat       685   803      750       586
## Fish             147   160      122        93
## Fats_and_oils    193   235      184       209
## Sugars           156   175      147       139
## Fresh_potatoes   720   874      566      1033
## Fresh_Veg        253   265      171       143
## Other_Veg        488   570      418       355
## Processed_potatoes 198   203      220       187
## Processed_Veg    360   365      337       334
## Fresh_fruit     1102  1137      957       674
## Cereals          1472  1582     1462      1494
## Beverages         57    73        53        47
## Soft_drinks     1374  1256     1572      1506
## Alcoholic_drinks  375   475      458       135
## Confectionery     54    64        62        41
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

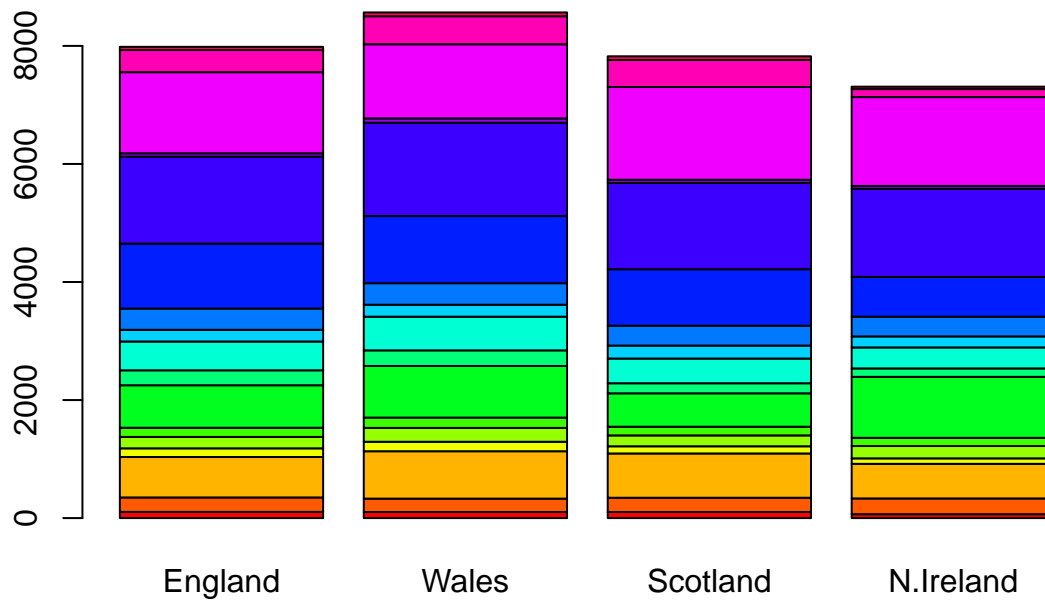
If you run the first code block multiple times you will keep losing columns. Using the `row.names` argument is more effective because you will prevent loss of data.

```
barplot(as.matrix(x), beside = T, col = rainbow(nrow(x)))
```



Q3. Changing what optional argument in the above `barplot()` function results in the following plot?


```
barplot(as.matrix(x), col = rainbow(nrow(x)))
```



Remove the 'beside = T' argument

Q5 (misabeled is Q4) Generating all pairwise plots. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col = rainbow(10), pch = 16)
```

> Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The points that are not on the diagonal (the blue and orange points) are different than the other countries.

PCA to the rescue

The main function in base R is 'prcomp()' This want's the transpose of our data

```
pca <- prcomp(t(x))
summary(pca)
```

```
## Importance of components:
##
##          PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

```
attributes(pca)
```

```
## $names
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
## $class
## [1] "prcomp"
```

```

plot(pca$x[,1], pca$x[,2], xlab = "PC1", ylab = "PC2", xlim = c(-270, 500))

## Warning in plot.window(...): "xlim" is not a graphical parameter

## Warning in plot.xy(xy, type, ...): "xlim" is not a graphical parameter

## Warning in axis(side = side, at = at, labels = labels, ...): "xlim" is not a
## graphical parameter


## Warning in axis(side = side, at = at, labels = labels, ...): "xlim" is not a
## graphical parameter

## Warning in box(...): "xlim" is not a graphical parameter

## Warning in title(...): "xlim" is not a graphical parameter

text(pca$x[,1], pca$x[,2], colnames(x))

```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```

plot(pca$x[,1], pca$x[,2], xlab = "PC1", ylab = "PC2", xlim = c(-270, 500))

## Warning in plot.window(...): "xlim" is not a graphical parameter

## Warning in plot.xy(xy, type, ...): "xlim" is not a graphical parameter

## Warning in axis(side = side, at = at, labels = labels, ...): "xlim" is not a
## graphical parameter


## Warning in axis(side = side, at = at, labels = labels, ...): "xlim" is not a
## graphical parameter

## Warning in box(...): "xlim" is not a graphical parameter

## Warning in title(...): "xlim" is not a graphical parameter

text(pca$x[,1], pca$x[,2], colnames(x), col = rainbow(4))

```



```

v <- round(pca$sdev^2/sum(pca$sdev^2)*100)
v


```

```
## [1] 67 29 4 0
```


```
z <- summary(pca)
z$importance
```

```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.15019 212.74780 73.87622 4.188568e-14
## Proportion of Variance 0.67444 0.29052 0.03503 0.000000e+00
## Cumulative Proportion 0.67444 0.96497 1.00000 1.000000e+00
```

```
barplot(v, xlab = "Principal Component", ylab = "Percent Variation")
```

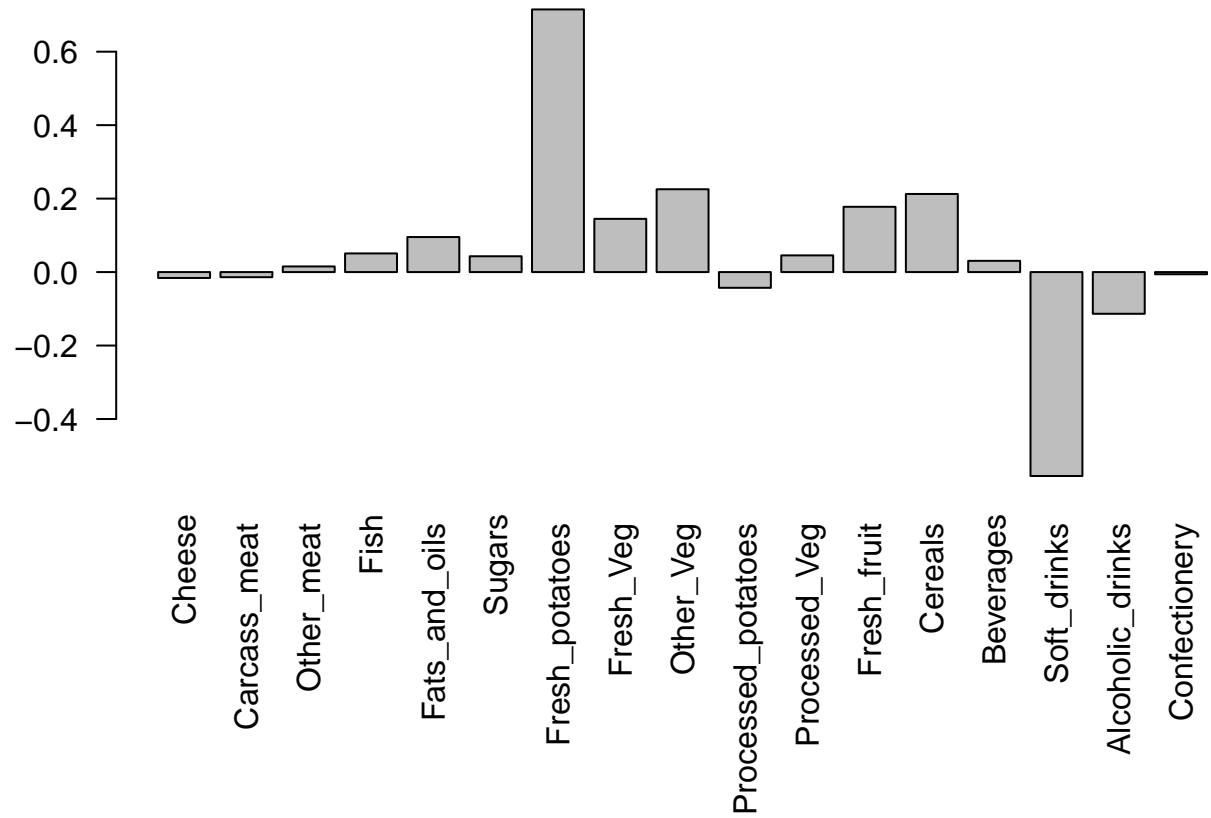


```
par(mar = c(10, 3, 0.35, 0))
barplot(pca$rotation[,1], las = 2)
```



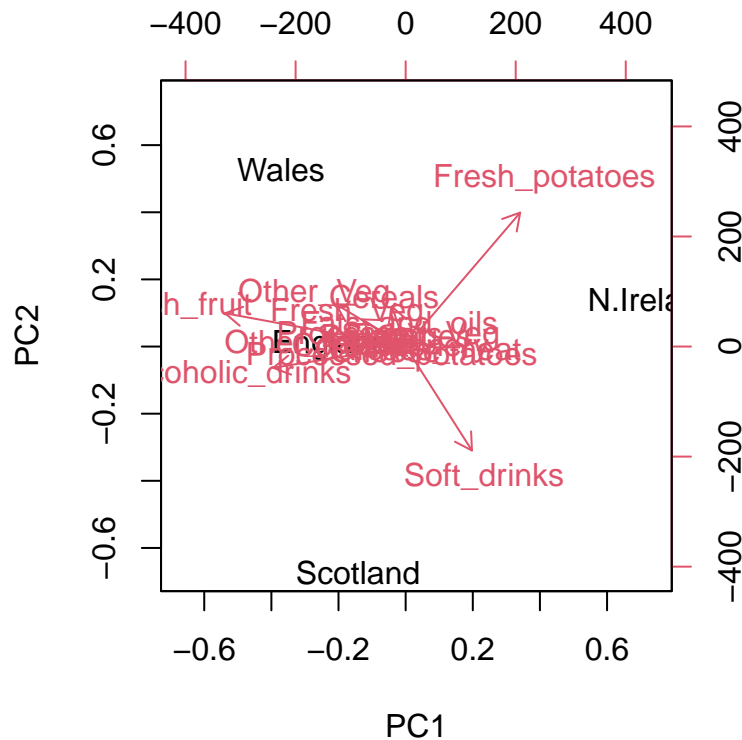
Q9. Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar = c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las =2)
```



PC1 reduces the data down into one dimension that covers about 67 percent of the data. PC2 covers is another dimension that covers another 29 percent of the data.

```
biplot(pca)
```



PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names = 1)
head(rna.data)
```

##	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
## gene1	439	458	408	429	420	90	88	86	90	93
## gene2	219	200	204	210	187	427	423	434	433	426
## gene3	1006	989	1030	1017	973	252	237	238	226	210
## gene4	783	792	829	856	760	849	856	835	885	894
## gene5	181	249	204	244	225	277	305	272	270	279
## gene6	460	502	491	491	493	612	594	577	618	638

Q10 How many genes and samples are in this data set?

```
dim(rna.data)
```


```
## [1] 100 10
```

Again we have to take the transpose of our data

```
pca <- prcomp(t(rna.data), scale = TRUE)
```

```
#Simple unpolished plot of pc1 and pc2
```

```
plot(pca$x[,1], pca$x[,2], xlab = "PC1", ylab = "PC2")
```



```
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
##              PC8      PC9      PC10
## Standard deviation  0.62065 0.60342 3.348e-15
## Proportion of Variance 0.00385 0.00364 0.000e+00
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

```
plot(pca, main = "Quick scree plot")
```


Quick scree plot




```
pca.var <- pca$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
## [1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

```
barplot(pca.var.per, main = "Scree plot", names.arg = paste0("PC", 1:10), xlab = "Principal Component",
```


Scree plot



```
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```

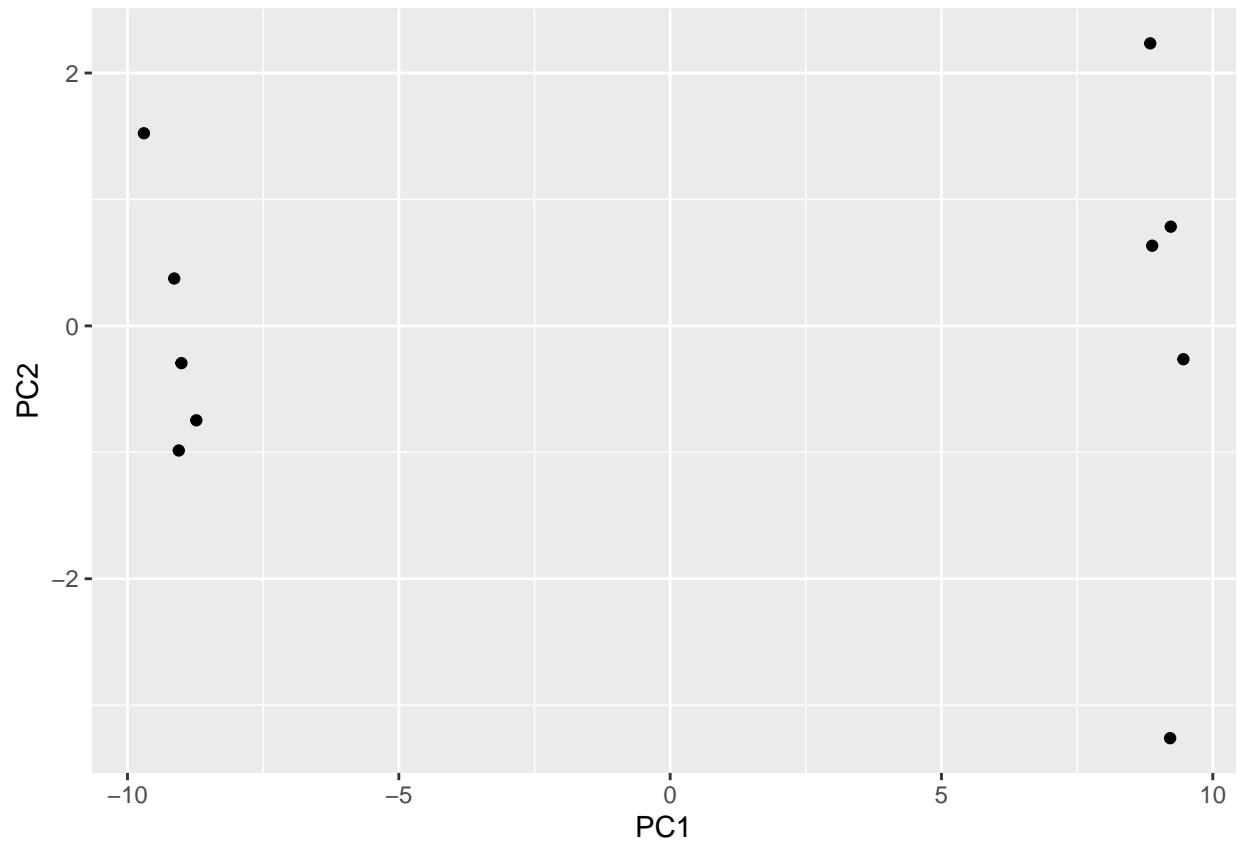


Use ggplot

```
library(ggplot2)

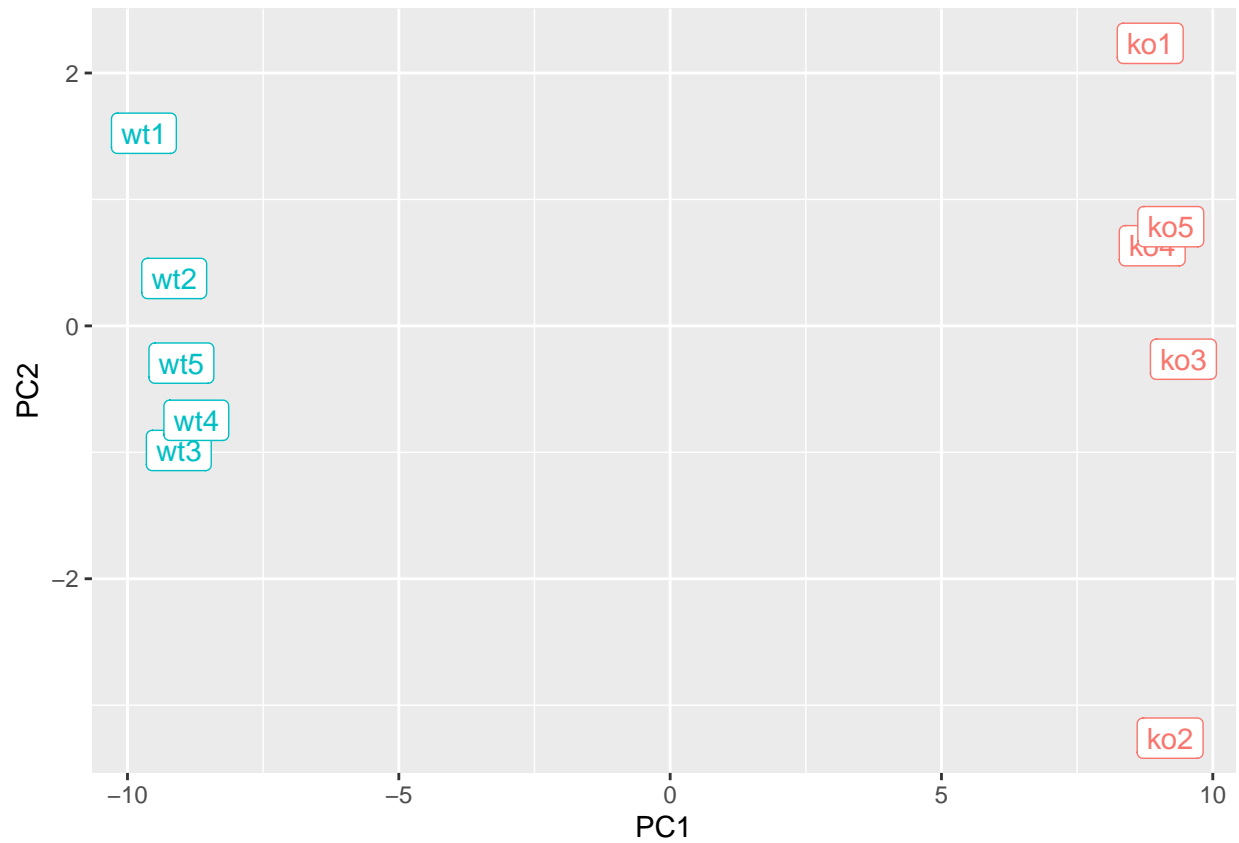
df <- as.data.frame(pca$x)

# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

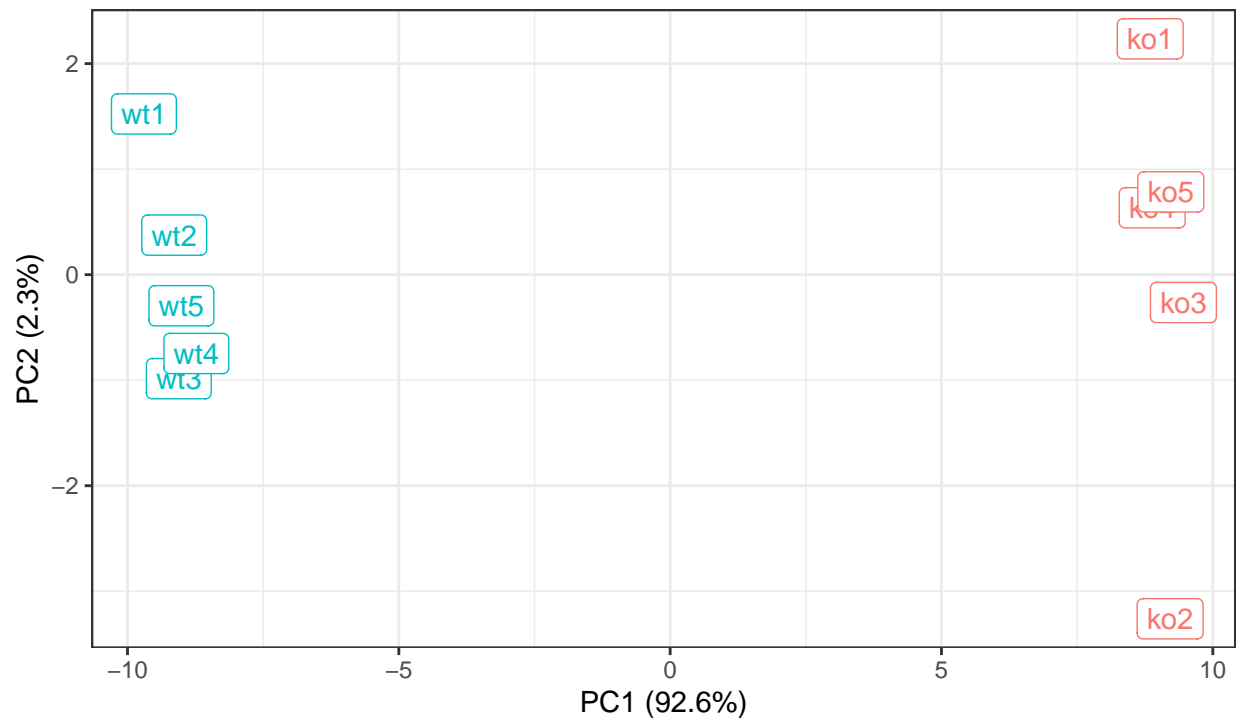
p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="BIMM143 example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



BIMM143 example data

```
loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes
```

```
## [1] "gene100" "gene66" "gene45" "gene68" "gene98" "gene60" "gene21"
## [8] "gene56" "gene10" "gene90"
```